

Chaitali Ranalkar

✉ chaitals1412@gmail.com | ☎ +91 98922 68972 | 📍 Pune, India

🌐 chaitals1412.github.io/research | 🐙 github.com/chaitals1412 | 🔗 linkedin.com/in/chaitali-ranalkar-1410cr

RESEARCH INTERESTS

- Data systems for AI: how LLMs can reliably clean, query and analyze messy real-world tabular data
- Reliable, uncertainty-aware models: knowing when a model's predictions and explanations can be trusted
- Learning from tabular and time-series data when standard modeling assumptions break down

EDUCATION

St. Xavier's College (Autonomous), Mumbai

Aug 2023 – Jun 2025

M.Sc. in Big Data Analytics

CGPA: 9.35 / 10

Relevant coursework: machine learning, time-series analysis

B. N. Bandodkar College

Jun 2020 – Jun 2023

B.Sc. in Mathematics

CGPA: 9.25 / 10

Relevant coursework: linear algebra, probability

RELEVANT PROJECTS

Quick-Commerce Operations Gap Analysis — Independent Project · Python, BigQuery, Power BI

Sep 2026

- Modelled a simulated 10-minute grocery-delivery platform (dark stores, inventory, delivery times, refunds, support tickets) to find where customer demand and revenue leak across the order lifecycle.
- Found that 54.3% of 150,000 searches ended with the product unavailable, so the largest loss occurs *before* an order is placed rather than during delivery; built Power BI views for platform health, leakage by area, and refund and delay follow-up.
- Traced an outlier store (6.1% on-time vs. a 42% network average) through a documented 26-query BigQuery root-cause chain, ruling out four hypotheses (peak-hour load, distance, workload, picking) and flagging it for on-site review instead of guessing a cause.

Customer Segmentation with RFM Features and K-Means — M.Sc. Research Project, Dept. of IT

Oct 2024

- Engineered log-transformed, min-max normalized Recency-Frequency-Monetary features from raw transactions; K-means (elbow at $K = 3$) produced three interpretable segments, benchmarked against rule-based RFM quintile scoring.
- Showed the two methods agree at the extremes but diverge in the middle, and analyzed where K-means' assumptions break (non-spherical clusters, outlier-driven centroids), motivating GMM, DBSCAN and hierarchical alternatives.

Predicting Customer Churn: Five Classifier Families Compared — M.Sc. Coursework

Mar 2024

- On 10,000 bank customers, a random forest tuned with 5-fold cross-validated search reached 88.0% accuracy, F1 0.83 and ROC-AUC 0.94, against 75.7% for logistic regression, which points to a non-linear decision boundary.
- Identified age as the dominant predictor (importance ≈ 0.28 , about twice the next feature), then examined the limits of that finding: impurity-importance bias, class imbalance ($\sim 20\%$ churners) and uncalibrated probabilities.

Forecasting ETH-USD: LSTM vs. Statistical Baselines — TensorFlow/Keras, statsmodels, Prophet

2024

- Modelled 1,827 daily closes of a non-stationary series (ADF $p = 0.34$); the LSTM reached a test RMSE of 135.8 vs. 1,031 (SARIMA), 1,040 (Holt-Winters) and 1,594 (Prophet), capturing turning points rather than only the trend.

PROFESSIONAL EXPERIENCE

Senior Executive, KKC & Associates LLP **Jul 2025 – Present | Pune**

- Built an IT general controls (ITGC) automation tool (Python, Streamlit) with a fully on-premises NLP pipeline for document summarization and policy search, so confidential audit data never leaves the firm; manual review effort fell by 50%.
- Developed Python tools that extract and consolidate data from many unstructured PDFs into clean, structured datasets for analysis.
- Automated ITGC testing and GST tax-return (R1/3B) reconciliation (extraction, validation, mismatch reporting) for 100+ reconciliations per cycle, cutting manual effort by 40%.
- Designed a Power BI/DAX project and partner management dashboard with row-level security, year-over-year cost comparison, and billing and utilization breakdowns by audit service line.

Data Analyst Intern, KKC & Associates LLP **Jan 2025 – Jun 2025 | Mumbai**

- Analyzed audit data with SQL, Power BI, Excel and CaseWare IDEA; built analytics covering financial performance, partner utilization, billing and project tracking.
- Prototyped the Python NLP summarizer for long policy and audit documents that was later developed into the ITGC tool.

POSTER PRESENTATION

Applications of Fourier Series — *B.Sc. Mathematics · Third Prize* **2020 – 2023**

TECHNICAL SKILLS

Programming	Python, R, SQL, DAX
Machine Learning & Statistics	scikit-learn, XGBoost, TensorFlow/Keras, statsmodels, Prophet; clustering, ensemble methods, time-series forecasting, NLP
Generative AI	Prompt engineering, LLM-assisted SQL generation
Data & Cloud Platforms	Google BigQuery, Google Cloud, Google Colab
Analytics & Tools	Power BI, Tableau, Streamlit, Excel, CaseWare IDEA, Power Automate

CERTIFICATIONS

Generative AI SQL Database Specialist — Vanderbilt University (Coursera)	Jul 2026
Introduction to Cloud Computing — Simplilearn	Jan 2024
Machine Learning — Great Learning Academy	Dec 2023